

159.341 Programming Languages, Algorithms & Concurrency

C++ `std::threads`

Daniel Playne
`d.p.playne@massey.ac.nz`

C++ `std::threads`

Threads were introduced to the C++ language with the release of the C++11 standard. The goal was tough, to provide a system for threads that is as efficient as a natively provided OS library but is portable between different operating systems.

The major advantage of threads being supported by the language is that only one implementation of a program needs to be developed and can simply be compiled for different operating systems.

C++ std::threads

Much of the C++ threading functionality is available in the header:

```
#include <thread>
```

The main object to look at is the `std::thread` which is a ***thread*** object or ***worker*** thread. It can be used to create a new thread and start executing a function.

C++ std::threads

Example 1:

```
#include<iostream>
#include<thread>

void thread_function() {
    std::cout << "Hello from thread" << std::endl;
}

int main() {
    std::thread t(thread_function);
    std::cout << "Hello from main" << std::endl;
    t.join();
}
```

Output (maybe):

```
Hello from main
Hello from thread
```

C++ `std::threads`

When an `std::thread` is created, an associated thread of control will be created and begin execution immediately (pending scheduling by the OS).

In general, each `std::thread` object is associated with one thread of execution and a thread of execution may be associated with at most one `std::thread` object.

C++ `std::threads`

There are four cases when an `std::thread` will not be associated with a thread of execution.

- Default construction
- After *move* construction (from)
- `join()` was called
- `detach()` was called

If an `std::thread` object is destructed while it is associated with a thread of execution, `std::terminate` will be called.

C++ `std::threads`

The most common usage is shown in the previous example where a thread is created and given a function to execute. The thread will run this function concurrently with the main thread still running the main function.

When the main thread wishes to end, it will call `join()` on the thread object `t` which will block until the thread of execution associated with `t` has ended.

If `join` is not called then `t` will still be associated with a thread of execution when it is destructed and cause an error.

C++ std::threads

Example 1 (no join):

```
#include<iostream>
#include<thread>

void thread_function() {
    std::cout << "Hello from thread" << std::endl;
}

int main() {
    std::thread t(thread_function);
    std::cout << "Hello from main" << std::endl;
}
```

Output (maybe):

```
Hello from main
libc++abi.dylib: terminating
Abort trap: 6
```


C++ `std::threads`

In most cases you will `join` your threads at the end of the program. However, you may encounter a case where you must destruct the `std::thread` object but allow the thread of execution to continue.

In this case you can `detach` a thread which will allow it to execute independently from the `std::thread` object used to create it.

C++ std::threads

Example 2:

```
#include<iostream>
#include<thread>

void thread_function() {
    std::cout << "Hello from thread" << std::endl;
}

int main() {
    {
        std::thread t(thread_function);
        t.detach();
    }
    std::cout << "Hello from main" << std::endl;
    std::this_thread::sleep_for (std::chrono::seconds(1));
}
```

Output (maybe):

```
Hello from main
Hello from thread
```

C++ `std::threads`

Once an `std::thread` has been detached it will no longer be associated with a thread of control and `join` cannot be called.

Detaching a thread does not mean it will be able to continue to execute if the main thread ends.

Usually, when the main function terminates the enclosing process will be destroyed and any associated threads terminated (without calling `std::terminate` however).

C++ std::threads

The other cases where an `std::thread` object is not attached to a thread of control are default construction and a move construction or assignment.

Example 3:

```
int main() {
    std::cout << "Hello from main" << std::endl;

    std::thread t1;                // t1 has no thread
    std::thread t2(thread_function); // t2 has thread
    std::thread t3(std::move(t2));  // thread moved to t3
    t1 = std::move(t3);            // thread moved to t1

    std::this_thread::sleep_for (std::chrono::seconds(4));

    t1.join();
    std::cout << "Goodbye from main" << std::endl;
}
```

C++ std::threads - Parameters

C++ syntax makes passing parameters to threads significantly simpler than a C API such as pthreads.

Example 4:

```
void thread_function(int start, int end) {  
    for(int i = start; i < end; i++) {  
        std::cout << i << std::endl;  
    }  
}  
  
int main() {  
    std::thread t1(thread_function, 0, 5);  
    std::thread t2(thread_function, 5, 10);  
    std::thread t3(thread_function, 10, 15);  
  
    t1.join(); t2.join(); t3.join();  
}
```

What will the output of this program be?

Parameters for C++ `std::threads`

Possible output:

05		0510		05
110		11		10
		12		11
116		13		12
7		14		13
2				14
3	or	1	or	1
4		2		2
8		3		6
9		4		7
		6		8
12		7		9
13		8		3
14		9		4

C++ std::threads - Parameters

Parameters can be passed by reference but must use `std::ref` to create a `std::reference_wrapper`.

Example 5:

```
void thread_function(int start, int end, long long int &total) {
    for(int i = start; i < end; ++i) {
        total += i;
    }
}

int main() {
    long long int total = 0;
    std::thread t0(thread_function, 0, 500, std::ref(total));
    std::thread t1(thread_function, 500, 1000, std::ref(total));
    t0.join();
    t1.join();
    std::cout << total << std::endl;
}
```

What will the output of this program be?

C++ `std::threads` - Parameters

Output (maybe):

378845

or

499500

or

445626

Lambda Functions with C++ `std::threads`

Threads can also be created to execute lambda functions:

Example 6:

```
int main() {  
    std::vector<std::thread> threads;  
    for(int i = 0; i < 5; i++) {  
        threads.push_back(std::thread([i] {  
            std::cout << i << std::endl;  
        }));  
    }  
    for(std::thread &t : threads) {  
        t.join();  
    }  
}
```

Output (maybe):

01
4

23

To provide mutual exclusion, C++ provides mutex objects.

There are four basic types provided:

- `mutex`
- `timed_mutex`
- `recursive_mutex`
- `recursive_timed_mutex`

The `mutex` primitive represents a basic mutex and has three main functions

- `lock()` - locks the mutex and blocks the thread if mutex is not available.
- `try_lock()` - tries to lock the mutex and returns `false` if the mutex was not obtained.
- `unlock()` - releases the lock on the mutex.

This mutex provides exclusive, non-recursive semantics:

- A thread *holds* the mutex from the time it successfully calls `lock` or `try_lock` and until it calls `unlock`.
- When a thread *holds* a mutex, any other thread will block (`lock`) or return false (`try_lock`) if they attempt to claim it.
- A thread must not already *hold* a mutex before calling `lock` or `try_lock`.

Example 7:

```
std::mutex mtx;

void thread_function(int id) {
    mtx.lock();
    std::cout << id << " counting to 3" << std::endl;
    for(int i = 0; i < 3; ++i) {
        std::cout << i << " ";
    }
    std::cout << std::endl;
    mtx.unlock();
}

int main() {
    std::vector<std::thread> threads;
    for(int i = 0; i < 5; i++) {
        threads.push_back(std::thread(thread_function, i));
    }
    for(std::thread &t : threads) {
        t.join();
    }
}
```

Output (maybe):

```
Thread: 1 counting to 3
0 1 2
Thread: 2 counting to 3
0 1 2
Thread: 0 counting to 3
0 1 2
Thread: 3 counting to 3
0 1 2
Thread: 4 counting to 3
0 1 2
```

C++ - `std::timed_mutex`

A thread may wish to obtain a lock on a mutex, however there may be cases when that mutex may be held by another thread for a significant period of time.

In such situations, being able to specify a timeout period can be useful to stop threads waiting too long to obtain a lock on a mutex.

The `timed_mutex` primitive extends `mutex` by including functions that allow a thread to include a maximum amount of time to wait when trying to claim a mutex.

C++ - `std::timed_mutex`

In addition to the functions provided by `mutex` it also provides:

- `try_lock_for()` - try to lock the mutex and return `false` if the mutex has been unavailable for a specified duration.
- `try_lock_until()` - try to lock the mutex and return `false` if the mutex has been unavailable until a specified timepoint has been reached.

C++ - std::timed_mutex

Example 8:

```
std::timed_mutex tmtx;

void thread_function(int id) {
    if(tmtx.try_lock_for(std::chrono::milliseconds(1000))) {
        std::this_thread::sleep_for(std::chrono::milliseconds(330));
        std::cout << "Thread " << id << " done." << std::endl;
        tmtx.unlock();
    } else {
        std::cout << "Thread " << id << " gave up." << std::endl;
    }
}

int main() {
    std::vector<std::thread> threads;
    for(int i = 0; i < 5; i++) {
        threads.push_back(std::thread(thread_function, i));
    }
    for(std::thread &t : threads) {
        t.join();
    }
}
```

C++ - std::timed_mutex

Output (maybe):

```
Thread 1 done.  
Thread 3 done.  
Thread 2 done.  
Thread 4 gave up.  
Thread 0 gave up.
```

or

```
Thread 1 done.  
Thread 0 done.  
Thread 3 done.  
Thread 4 gave up.  
Thread 2 done.
```

C++ - `std::recursive_mutex`

Sometimes (often in a recursive function) it is not always straightforward to determine whether a thread has already obtained a lock on a mutex or not.

For this case, C++ provides another type of mutex called `std::recursive_mutex`.

This mutex may be locked multiple times (by the same thread) and is considered to be unlocked when the matching number of unlock calls have been made.

C++ - `std::recursive_mutex`

This mutex provides exclusive, recursive semantics:

- A thread *holds* the recursive mutex from the time it successfully calls `lock` or `try_lock`. Once held, the thread may make additional calls to `lock` and `try_lock` and will continue to hold the mutex until the matching number of calls to `unlock` have been made.
- When a recursive mutex is held by a thread, calls to `lock` or `try_lock` made by any other thread will block or return false.
- The maximum number of times a `recursive_mutex` may be locked by a thread is unspecified but will throw `std::system_error (lock)` or return false `try_lock` when reached.

C++ - std::recursive_mutex

Example 9:

```
std::recursive_mutex rmtx;

void thread_function(int id) {
    rmtx.lock();
    std::cout << id << " ";
    if(id > 0) {
        thread_function(id - 1);
    } else {
        std::cout << std::endl;
    }
    rmtx.unlock();
}

int main() {
    std::vector<std::thread> threads;
    for(int i = 0; i < 5; i++) {
        threads.push_back(std::thread(thread_function, i));
    }
    for(std::thread &t : threads) {
        t.join();
    }
}
```

C++ - std::recursive_mutex

Output (maybe):

```
1 0
0
3 2 1 0
4 3 2 1 0
2 1 0
```

or

```
1 0
2 1 0
3 2 1 0
4 3 2 1 0
0
```

C++ - std::lock

Another mechanism for locking a mutex is to use the function `std::lock` which allows multiple mutexes to be locked at the same time while avoiding deadlock (order does not matter).

```
std::mutex mtx1, mtx2;

void thread_function_a() {
    std::lock(mtx1, mtx2);
    std::cout << "Thread a" << std::endl;
    mtx1.unlock(); mtx2.unlock();
}

void thread_function_b() {
    std::lock(mtx2, mtx1);
    std::cout << "Thread b" << std::endl;
    mtx1.unlock(); mtx2.unlock();
}

int main() {
    std::thread t_a(thread_function_a);
    std::thread t_b(thread_function_b);
    t_a.join();
    t_b.join();
}
```

C++ - `std::lock`

The definition of `std::lock` states that all mutexes will be locked via a sequence of calls to `lock`, `try_lock` and `unlock` and will not result in deadlock. The exact method for avoiding deadlock is not specified (to avoid constraining implementations).

The definition must be interpreted with care, using `std::lock` to lock mutexes does not mean that deadlock cannot occur globally throughout the program.

It means that within a single call to `std::lock` the thread will not hold the lock on any mutex while waiting for another to become available.

C++ - std::lock

Deadlock can still occur when multiple calls to std::lock are made.

```
std::mutex mtx1, mtx2, mtx3, mtx4;

void thread_function_a() {
    cout << "Wait - 1,2 (a)" << endl;
    std::lock(mtx1, mtx2);
    cout << "Wait - 3,4 (a)" << endl;
    std::lock(mtx3, mtx4);
    mtx1.unlock(); mtx2.unlock();
    mtx3.unlock(); mtx4.unlock();
}

int main() {
    std::thread t_a(thread_function_a);
    std::thread t_b(thread_function_b);
    t_a.join();
    t_b.join();
}
```

```
void thread_function_b() {
    cout << "Wait - 3,4 (b)" << endl;
    std::lock(mtx3, mtx4);
    cout << "Wait - 1,2 (b)" << endl;
    std::lock(mtx1, mtx2);
    mtx1.unlock(); mtx2.unlock();
    mtx3.unlock(); mtx4.unlock();
}
```

C++ `std::threads` - **Mutexes**

Mutexes are generally not intended to be used directly and more commonly utilised through `lock_guard` or `std::unique_lock`.

A `lock_guard` is an RAI (Resource Allocation Is Initialisation) mechanism for taking ownership of a mutex for the duration of a scoped block.

A `unique_lock` is a general-purpose mutex ownership wrapper that supports deferred locking, time-constrained attempts at locking, recursive locking, transfer of ownership and use with condition variables.

Summary

C++ threads:

- `thread`
- `mutex`
- `lock`