

159.341 Programming Languages, Algorithms & Concurrency

C++ `std::threads`

Daniel Playne
`d.p.playne@massey.ac.nz`

Reminder

In the last lesson we looked at some of the basics of C++ threading.

- `thread`
- `mutex`
- `timed_mutex`
- `recursive_mutex`
- `recursive_timed_mutex`
- `lock`

C++ `std::threads` - **Mutexes**

Mutexes are generally not intended to be used directly and more commonly utilised through `lock_guard` or `std::unique_lock`.

A `lock_guard` is an RAI (Resource Allocation Is Initialisation) mechanism for taking ownership of a mutex for the duration of a scoped block.

A `unique_lock` is a general-purpose mutex ownership wrapper that supports deferred locking, time-constrained attempts at locking, recursive locking, transfer of ownership and use with condition variables.

C++ std::threads - lock_guard

When a `lock_guard` object is created, it will attempt to take ownership of the given mutex. When the `lock_guard` leaves scope and is destructed, it will release the mutex.

```
std::mutex mtx;

void thread_function(int id) {
    std::this_thread::sleep_for(std::chrono::seconds(1));
    std::lock_guard<std::mutex> lck(mtx);
    std::cout << "Thread: " << id << std::endl;
}

int main() {
    std::vector<std::thread> threads;
    for(int i = 0; i < 5; i++) {
        threads.push_back(std::thread(thread_function, i));
    }
    for(std::thread &t : threads) {
        t.join();
    }
}
```

C++ `std::threads` - `lock_guard`

One of the main advantages of using a `lock_guard` is that they ensure mutexes are properly unlocked when an exception occurs.

If a `lock_guard` is created inside a try block and an exception occurs. The mutex will be properly released when the exception-handling piece of code is executed and the `lock_guard` leaves scope.

C++ std::threads - lock_guard

```
std::mutex mtx;

void print_maybe(int i) {
    int r = rand();
    if(r % 2 == 0) {
        std::cout << i << std::endl;
    } else {
        throw "Unlucky";
    }
}

void thread_function(int id) {
    try {
        std::lock_guard<std::mutex> lock(mtx);
        print_maybe(id);
    } catch(const char *s) {
        // Handle exception
    }
}

int main() {
    std::vector<std::thread> threads;
    for(int i = 0; i < 5; i++) {
        threads.push_back(std::thread(thread_function, i));
    }
    for(std::thread &t : threads) {
        t.join();
    }
}
```

C++ `std::lock_guard`

There is another constructor for `lock_guard`

```
lock_guard(mutex_type &m, std::adopt_lock_t t)
```

This constructor will create a `lock_guard` and adopt the ownership of the mutex without trying to lock it.

Adopting the lock of a mutex should only ever be used if the thread already owns the mutex (behaviour is undefined otherwise).

C++ - std::lock_guard

The `adopt_lock` constructor allows `lock_guard`s to be used in conjunction with `std::lock`. First the necessary mutexes are locked, then `lock_guard`s are constructed that adopt those locks.

```
std::mutex mtx1, mtx2, mtx3;

void thread_function_a() {
    std::lock(mtx1, mtx2, mtx3);
    std::lock_guard<std::mutex> lck1(mtx1, std::adopt_lock);
    std::lock_guard<std::mutex> lck2(mtx2, std::adopt_lock);
    std::lock_guard<std::mutex> lck3(mtx3, std::adopt_lock);

    std::cout << "Thread a" << std::endl;
}

void thread_function_b() {
    std::lock(mtx3, mtx2, mtx1);
    std::lock_guard<std::mutex> lck1(mtx1, std::adopt_lock);
    std::lock_guard<std::mutex> lck2(mtx2, std::adopt_lock);
    std::lock_guard<std::mutex> lck3(mtx3, std::adopt_lock);

    std::cout << "Thread b" << std::endl;
}
```


C++ `unique_lock`

C++ also has `unique_lock` which supports *move semantics* to allow a `unique_lock` to be *moved*.

This feature can be useful when a thread that currently holds a mutex wishes to pass ownership of that mutex to another thread.

The original thread can *move* the `unique_lock` to the other thread and ensure that no other thread obtains ownership of it in between.

C++ unique_lock

Moving a unique_lock to another thread:

```
std::mutex mtx1;

void thread_function(int id, std::unique_lock<std::mutex> &&lck) {
    std::this_thread::sleep_for(std::chrono::seconds(1));
    std::cout << "Thread " << id << " running." << std::endl;
    std::this_thread::sleep_for(std::chrono::seconds(1));
}

int main() {
    // Create Threads
    std::vector<std::thread> threads;
    for(int i = 0; i < 5; i++) {
        // Get unique lock
        std::unique_lock<std::mutex> lck(mtx1);

        // Create Thread (move unique lock to thread)
        std::cout << "Lanching thread " << i << "." << std::endl;
        threads.push_back(std::thread(thread_function, i, std::move(lck)));
    }

    // Join Threads
    for(std::thread &t : threads) {
        t.join();
    }
}
```

C++ `condition_variable`

The C++ threading API (before C++20) does not support *semaphores*.

Instead a `condition_variable` can be used to block a thread until *notified* by another thread.

While a thread is waiting on a `condition_variable` it will be blocked and not scheduled on the CPU.

C++ condition_variable

`condition_variable` provides the following functions:

- `wait` - Wait until notified
- `wait_for` - Wait until notified or timeout
- `wait_until` - Wait until notified or time point is reached
- `notify_one` - Notify one thread waiting on condition variable.
- `notify_all` - Notify all threads waiting on condition variable.

C++ condition_variable

Example:

```
std::mutex mtx;
std::condition_variable cond;
bool ready = false;

void thread_function(int id) {
    std::unique_lock<std::mutex> lock(mtx);
    std::cout << "Thread " << id << " waiting." << std::endl;
    while(!ready) cond.wait(lock);
    std::cout << "Thread " << id << " completed." << std::endl;
}

int main() {
    std::vector<std::thread> threads;
    for(int i = 0; i < 5; ++i) {
        threads.push_back(std::thread(thread_function, i));
    }
    std::this_thread::sleep_for (std::chrono::seconds(3));
    ready = true;
    cond.notify_all();
    for(std::thread &t : threads) {
        t.join();
    }
}
```

C++ condition_variable

In the following thread function, the thread obtains a lock on the mutex `mtx` which will be held until the call to `cond.wait(lock)`.

When the thread calls `wait` it will release the lock. When the thread is *notified* it will wake up and obtain a lock on `mtx` before continuing execution.

```
std::mutex mtx;
std::condition_variable cond;
bool ready = false;

void thread_function(int id) {
    std::unique_lock<std::mutex> lock(mtx);
    std::cout << "Thread " << id << " waiting." << std::endl;
    while(!ready) cond.wait(lock);
    ...
}
```

C++ condition_variable

Note the following line is used when *waiting* on a condition variable.

```
while(!ready) cond.wait(lock);
```

In general, threads are only woken up when another thread calls either `notify_one` or `notify_all`. However, some implementations may produce *spurious* calls to wake up threads.

It is the programmer's responsibility to ensure the correct conditions to resume execution have been met.

Example - Game of Life

To continue our Game of Life example we should look at how to implement our *phase-parallel* simulator.

We will need to explicitly create threads that will each update one section of the grid.

After each new generation is computed, the threads must synchronise with each other before continuing onto the next generation.

Example - Game of Life

First a sequential implementation

```
unsigned char *buffer[2];
buffer[0] = new unsigned char[N*N];
buffer[1] = new unsigned char[N*N];

int r = 0; // Read/Write Index
int w = 1;

for(int k = 0; k < N*N; ++k) {    // Initialise
    buffer[r][k] = rand() % 2;
}

for(int ig = 0; ig < G; ++ig) {    // For each Generation
    for(int iy = 0; iy < N; ++iy) {    // For each Row
        for(int ix = 0; ix < N; ++ix) {    // For each Column
            // Count neighbours
            int count = count_neighbours(buffer[r], ix, iy, N);

            // Update
            buffer[w][iy*N + ix] = gol_update(buffer[r][iy*N+ix], count);
        }
    }

    // Swap buffers
    r = !r;
    w = !w;
}
```

Example - Game of Life

Game of Life Functions:

```
// Count neighbours
int count_neighbours(unsigned char *buffer, int ix, int iy, int N) {
    // Neighbouring Indexes
    int ym1 = (iy == 0) ? N-1 : iy-1;
    int yp1 = (iy == N-1) ? 0 : iy+1;
    int xm1 = (ix == 0) ? N-1 : ix-1;
    int xp1 = (ix == N-1) ? 0 : ix+1;

    // Count neighbours
    return buffer[ym1 * N + xm1] + buffer[ym1 * N + ix] + buffer[ym1 * N + xp1] +
           buffer[iy * N + xm1] +                          buffer[iy * N + xp1] +
           buffer[yp1 * N + xm1] + buffer[yp1 * N + ix] + buffer[yp1 * N + xp1];
}

// Game of Life Rules
unsigned char gol_update(unsigned char status, int count) {
    if((status == 0) && (count == 3)) {
        return 1;
    } else if((status == 1) && (count == 2 || count == 3)) {
        return 1;
    } else {
        return 0;
    }
}
```

Example - Game of Life

Creating threads to update different sections of the grid is relatively straightforward.

There is still the question of how many threads to use, which will we avoid by making the user decide.

```
std::thread *threads = new std::thread[T];

for(int t = 0; t < T; ++t) {
    threads[t] = std::thread(gol, ...);
}
```

Example - Game of Life

```
// Thread Function
void gol(unsigned char *buffer[2], int N, int G, int id, int num_threads) {
    int r = 0; // Read/Write index
    int w = 1;

    int start = ( id      * N) / num_threads; // Start Row
    int end   = ((id+1) * N) / num_threads; // End Row

    for(int ig = 0; ig < G; ++ig) {          // For each Generation
        for(int iy = start; iy < end; ++iy) { // For each Row (for this thread)
            for(int ix = 0; ix < N; ++ix) {    // For each Column
                // Count neighbours
                int count = count_neighbours(buffer[r], ix, iy, N);

                // Update
                buffer[w][iy*N + ix] = gol_update(buffer[r][iy*N+ix], count);
            }
        }

        // Barrier
        barrier(ig+1, num_threads);

        // Swap buffers
        r = !r;
        w = !w;
    }
}
```

C++ condition_variable

To ensure that all the threads have updated their sections of the grid before continuing onto the next generation, we have used the following function:

```
barrier(ig+1, num_threads)
```

Barriers are only defined in the C++20 standard and may not be available in many compilers.

We will need to implement our own barrier function.

Example - Game of Life

```
std::condition_variable barrier_condition;
std::mutex barrier_mutex;

int barrier_count = 0;
int barrier_generation = 0;

void barrier(int g, int num_threads) {
    // Lock Mutex
    std::unique_lock<std::mutex> lock(barrier_mutex);

    // Increment barrier counter
    ++barrier_count;

    if(barrier_count == num_threads) {
        // Update barrier generation
        barrier_generation = g;

        // Reset barrier count
        barrier_count = 0;

        // Notify all waiting threads
        barrier_condition.notify_all();
    } else {
        // Wait
        barrier_condition.wait(lock, [=]() {
            return barrier_generation == g;
        });
    }
}
```

Summary

- `lock_guard`
- `unique_lock`
- `condition_variable`
- Game of Life implementation